

The Other Half of the Memory Wall: Serving 35B MoEs from SSD with Trained Routing Prediction

Yu Lin^{*}, Yiming Wang[†], Runyuan Cai[‡], Hanze Liu[§], Xiaodong Zeng[¶]
AutoArk

September 2026

Abstract

Mixture-of-experts (MoE) inference on consumer hardware is bounded by weight memory: a 35B-class model is 19.5 GB at 4-bit, and sparsity shrinks the compute per token, not the bytes that must be held. Naive offloading to SSD does not help on its own, because layer $N+1$'s experts must be chosen before layer N 's output exists, so the reads cannot start early enough to hide behind compute. We present EDGE0, a streaming MoE inference engine that closes the gap with a *prerouter*: a per-layer head predicts the next layer's routing one token ahead, and the prediction is consumed as the routing itself, so the staged expert set equals the routed set and nothing is dropped. An unmerged *recovery LoRA*, trained on the student path, pays back the quality lost to int4 quantization and routing replacement. On a single 24 GB machine, EDGE0 serves a 35B MoE at ≈ 20 tok/s inside 3 GiB of peak active memory, within a few points of its fp16 teacher on average across five public benchmarks. An 8B tier runs on the same framework, and the framework, checkpoints, and adapters are open source.

1 Introduction

In 1995 Wulf and McKee named the memory wall: processors improving faster than DRAM, with machines spending an ever-larger share of their cycles waiting for memory [32]. Thirty years later the wall confronts AI inference, and its arithmetic is unkind to the datacenter-free deployment of large models. Decode moves a few FLOPs per byte of weight it reads while hardware ridge points sit orders of magnitude higher; the constraint is bytes, and the bytes live in memory.

Memory holds two kinds of data that behave very differently. *State*, the KV cache, is dynamic and grows with every generated token. *Weights* are static, fixed the day training ends. The industry attacked the dynamic half: multi-head latent attention (MLA) compressed KV state by an order of magnitude [5], sparse attention capped it at a fixed horizon [7], linear state models removed the explicit cache altogether [13]. The bill that could not be renegotiated (tens of gigabytes of weights) was answered with one move: put it in a datacenter, where expert parallelism and sharded serving absorb it [25].

^{*}Corresponding author. E-mail: yu.lin@autoark.ai

[†]yiming.wang@autoark.ai

[‡]runyuan.cai@autoark.ai

[§]hanze.liu@autoark.ai

[¶]xiaodong.zeng@autoark.ai

The two standard local answers both fail at 35B scale. Quantization bottoms out at 4 bits; below that, error grows until the model is unusable, and no post-hoc technique recovers it [11, 21]. MoE sparsity helps on a different axis: a 35B MoE activates about 3B parameters per token, which shrinks what you *compute*, not what you *store*. Nineteen and a half gigabytes still have to be somewhere, and on a 24 GB desktop shared with an operating system, that somewhere is the machine’s entire memory, crowding out the OS and everything else the user is running.

EDGE0 changes where the weights live. Expert weights stay on SSD and are mmap-streamed into memory on demand, and peak memory is bounded by the active set instead of the parameter count. On-demand loading by itself is not fast enough: at every decode step, layer $N+1$ ’s expert selection depends on layer N ’s output, so a naive streaming engine stalls on disk latency once per layer per token. EDGE0 removes the stall with a trained *prerouter*: a small per-layer head that predicts layer $N+1$ ’s routing from layer N ’s state at the previous token, so expert reads overlap the forward pass. The prediction is consumed as the routing itself: the staged expert set and the routed set are identical by construction, and the approximation that other pre-gating schemes absorb at inference time through fallback loads and dropped tokens is instead paid once, in training, and recovered there.

Approximate routing and 4-bit quantization both cost quality. EDGE0 trains a *recovery LoRA* on top: the int4 base is frozen, a low-rank adapter is distilled from the fp16 teacher under the student routing path, and the adapter is served *unmerged* as a parallel delta. Merging and re-quantizing to 4-bit erases most of the adapter’s effect (§3.3).

Contributions:

1. **SSD as the weight tier.** An expert executor that streams MoE expert weights from disk on demand, with routing math pinned bit-identical to the vendored models and every executor path verified element-wise against the dequantized reference, plus a slot mechanism that removes the per-step stack-rebuild tax (§3.1).
2. **Prediction as routing.** A cross-token prerouter whose prediction *is* the routing at decode, so staged decode drops nothing and expert reads overlap compute, worth +80 to +84% decode on a machine where the checkpoint does not fit, together with the training recipe that makes a language model work under replaced routing (§3.2, §4, §5.3).
3. **Unmerged recovery LoRA.** Merging an adapter into an int4 base and re-quantizing destroys most of its effect; we show that serving it as a parallel delta avoids this at negligible cost (§3.3).
4. **Two open tiers, end to end.** A 35B and an 8B model released as checkpoint plus adapters that load together, each within a few points of its fp16 base on average (§5).

2 Background and Related Work

MoE routing. Production sparse MoEs route each token to K of E experts per layer [10, 18, 20, 26]. Two routing families dominate. Softmax-top-k (Qwen3.6): exact softmax over expert logits, top- k , renormalize. Sigmoid-group (DeepSeek-V3 [6], Ling [29]): sigmoid scores; groups ranked by the sum of their top two scores, best G groups survive; top- k within survivors; weights are the raw sigmoid values, renormalized and scaled. EDGE0 implements both verbatim and reuses the same two functions for prerouter training and inference (§3.2).

Pre-gating and the lead distance. Pre-gated MoE [17] selects the next block’s experts within the same token: the decision is produced after block N ’s attention and consumed before block $N+1$ ’s weights are fetched. That schedule does not survive contact with a streaming engine, and EDGE0 leads by a full token instead (§3.2, where we price the per-layer alternative).

KV-side compression. MLA [5], sparse attention [7], and linear-state models [13] compress dynamic state, and serving systems page it instead of compressing it [19]. They are orthogonal to EDGE0, and EDGE0 benefits from them: its 8B tier is an MLA + MoE hybrid whose compressed attention leaves more of the memory budget for the active expert set.

Quantization. GPTQ/AWQ-class methods [11, 21] made 4-bit the practical working point for post-training quantization; below 4 bits, post-hoc error grows fast enough that deployed local models rarely go lower. EDGE0 uses int4 affine group-64 expert weights and accepts the loss openly, then recovers most of it with distillation. QLoRA established that a 4-bit base carrying adapters can approach 16-bit fine-tuning quality [8]; we take the 4-bit base as given and ask what the adapter itself survives at serve time (§3.3). Closest in approach is expert-skip self-distillation [23], which shows a post-trained MoE tolerates halved expert counts when the model is trained for it; our routing-replacement training is the same observation applied to prediction-as-routing.

Offloading. Expert offloading is well studied: llama.cpp’s `--cpu-moe` [12] keeps MoE experts in CPU/system memory, PowerInfer [28] splits neurons into hot and cold sets across GPU and CPU, Mixtral-offloading [9] and MoE-Infinity [33] cache experts by popularity or by reuse distance, and FlexGen [27] extends the hierarchy to disk for weights and KV together. All of them move the footprint rather than shrinking it: the weights still occupy tens of gigabytes, and the ones that reach disk do so without knowing which experts the next step needs.

3 System Design

EDGE0 is a streaming MoE inference framework: all MLX code lives behind a backend facade (`core/nn/io/quant`), and core logic (model specs, prerouter, streaming pool, server) depends only on that facade, so additional backends implement the same surface. Models are described by a single `MoESpec` (expert count, K , routing family, quantization, weight layout, key templates); one generic streaming layer serves every model. LoRA and prerouter weights are safetensors files with provenance metadata, resolved from the model directory, and never merged into the base. Figure 1 shows the system architecture and the decode-time dataflow.

3.1 SSD Streaming Expert Layers

Problem. A 40-layer, 256-expert MoE holds 453 MB of expert weights per layer at int4—1.77 MB for each of its 256 experts—or 18 GB (16.9 GiB) of routed experts inside a 19.5 GB checkpoint whose remainder is attention, embeddings, the shared expert, and int4 group scales. That is three quarters of the RAM of a 24 GB machine, before KV cache, the operating system, and anything else the user runs. EDGE0 mmap’s the quantized expert safetensors and reads byte ranges on demand; the operating-system page cache carries hot data and peak memory tracks the *active* set, not the parameter count.




Figure 1: The EDGE0 system. Expert weights stay on SSD as int4 per-layer stacked safetensors and are mmap-streamed on demand into a bounded streaming pool (OS page cache, LRU, staged fixed-slot double buffer). A frozen int4 decoder stack computes with a parallel, unmerged recovery LoRA branch, $y = W_{\text{int4}}(x) + \frac{\alpha}{r} B A x$, combined at a sum node. The prerouter head owned by layer N predicts layer $N+1$ ’s routing one token ahead (red path, “prediction is the routing”), so the SSD reads that fill the staged slots overlap the forward pass; the decode timeline at the bottom shows the predicted read of token $t+1$ running concurrently with the compute of token t .

Executor paths. Four paths, from the correctness baseline to the prefill bulk path:

1. *exact*: deduplicated on-demand bundle build, stack, quantized gather. The correctness baseline for every other path.
2. *staged*: fixed-slot double buffering, the decode workhorse. Routing indices are mapped through a slot table with `take`; indices never leave the GPU. Repeated expert sets reuse cached graph nodes; persistent sticky-slot tensors are updated in place (`incr_stack`) so a step rewrites only the experts that changed instead of rebuilding the layer’s nine stacked tensors—gate, up, and down weights, each with its scales and biases—across all 40 layers.
3. *hot*: a fixed set of LRU-resident hot experts per layer, the same hot/cold split as PowerInfer and the popularity caches of Mixtral-offloading and MoE-Infinity [9, 28, 33]; hits take the stacked gather, misses fall to exact.
4. *whole-layer*: all experts of a layer loaded in one shot for prefill, which the checkpoint’s per-layer stacked layout makes nine direct reads—the same three quantized projections—rather than 256×9 builds; CPU load overlaps the previous layer’s GPU execution.

Math contract. Every path computes $\text{down}(\text{silu}(\text{gate}(x)) \cdot \text{up}(x))$ with the same quantized gather kernel, and the test suite compares each path element-wise against the dequantized reference (relative L2 < 1%; measured residual $\approx 0.24\%$ is the bf16-internal precision of the kernel itself). This contract is what allows the executor to switch paths freely per layer, per phase, without changing outputs.

Why offload can win outright. A fully-resident engine on this hardware is not the fast baseline it appears to be, because the model does not fit: the Mac mini M4 Pro has 24 GB, and a vanilla `mlx-lm` server with all 19.5 GB of 4-bit weights resident decodes at 3.9 tok/s occupying 18.2 GiB, while EDGE0’s $K=4$ profile decodes at 20.4 tok/s occupying 2.9 GiB.¹ The resident path holds 18.2 GiB that cannot be reclaimed, which leaves almost nothing of the 24 GB for the KV cache and the operating system; EDGE0 pays instead in graph building and tensor assembly, a cost of our streaming engine rather than of the weights, and the subject of §6. The oracle experiment makes that cost explicit: with perfect routing prediction and an unbounded cache, the only remaining cost is per-step tensor assembly, which `incr_stack` then attacks directly (Appendix B).

3.2 The Prerouter: Prediction Is the Routing

Motivation. Expert selection for layer $N+1$ depends on layer N ’s output, which does not exist yet when layer $N+1$ ’s loads would need to start. Waiting serializes disk latency into every step. The prerouter breaks the dependency with a double shift: the head owned by layer N runs at token t on layer N ’s post-attention norm output and predicts layer $N+1$ ’s routing at token $t+1$. Layer $N+1$ consumes the prediction made one token earlier; the SSD read that fills its slots overlaps the current forward pass (Figure 1).

The lead has to be a full token, not one layer. Choosing layer $N+1$ ’s experts right after layer N ’s attention, as Pre-gated MoE does [17], needs a per-layer synchronization and head evaluation that drain the GPU pipeline (30–100 ms per step in our engine, more than the load time it can hide, and every same-token variant we measured fell below a plain LRU baseline), and the next layer’s attention has not been computed when its experts must be chosen. One token of lead moves the head evaluation off the per-layer critical path: a single flush per step predicts every staged layer at once (32 on the 35B tier, 16 on the 8B tier), and the window it opens has to span a whole decode step, because the reads it hides are a step’s worth of disk traffic: at $K=4$ a layer’s experts are ≈ 7 MB, and the prerouter arm still spends 101.9 ms per step waiting on cold loads (§5.3).

Head. Per layer: $fc_1 \rightarrow \text{erf-gelu} \rightarrow fc_2$, plus a linear residual path ℓ that the training script warm-starts from the *next layer’s* router weight (its default is zero), so training begins from “apply the next router directly to this hidden state” and the MLP learns the correction. The input feature concatenates the hidden state with two top- k one-hots: the experts this layer actually routed to at this token, and those at the previous token ($\text{dim} = d_{\text{model}} + 2E$; for the 35B tier, $2048 + 2 \times 256 = 2560$, hidden width 512). Heads are fp16, the training export precision; running them in fp32 costs measurable time for no accuracy benefit. The 35B tier carries 33 heads (owners 6–38) and the 8B tier 16 (owners 7–22); predictions are consumed at layers 7–38 on the 35B tier and 8–23 on the 8B tier, so 32 of 40 and 16 of 24 layers stream from a prediction rather than from their own gate (the 35B head owned by layer 38 predicts layer 39’s routing, which is never staged, so it ships without a consumer). The 8B input is $1536 + 2 \times 128$.

Prediction-as-routing. At decode, MoE layers route by the prerouter’s logits instead of the router’s, through the same softmax-topk or sigmoid-group math as the original router (Appendix A). Because the routed set *is* the predicted set, the staged slots map exactly and there is nothing to drop: the coverage-vs-quality trade-off that pre-gated systems handle at runtime [17] is eliminated by construction. What the approximation costs is transferred to training, where it can be paid once (§4).

¹Same machine, same decode protocol, think-mode on.

Feature drift. A head’s input includes the routing its layer actually executed. At training that is the base router’s selection; at decode it is the head’s own prediction, because the prediction replaced the router. The heads are not retrained for that shift. The recovery LoRA is trained on the student path with prerouting in place, so it sees the deployed input distribution, and the quality measurements of §5 price the shift together with int4 and the routing approximation.

Two consumption profiles. Both tiers run prediction-as-routing staged decode: the 8B tier at $K=8$ with eight staged slots per layer, the 35B tier at $K=4$ with the incremental sticky-slot stack (§3.1). Prefill takes the whole-layer path on both tiers, where every expert of a layer is hit and there is nothing to predict, and the prerouter is exercised at decode. The same trained heads and the same stager abstraction serve both profiles; the framework exposes them as layer options.

3.3 Recovery LoRA: Unmerged by Design

The served model is (int4 base) + (routing replacement) + (LoRA), and the LoRA [15] exists to recover the first two. The sidecar is the same low-rank construction we used for multi-task and privacy-preserving serving [30, 31], and the same frozen-base-plus-sidecar pattern recently applied to generative-vision personalization [3]. It is trained on the student path (routing by prerouter) with cross-entropy against the teacher’s data, so the adapter compensates quantization and routing approximation jointly.

Deployment keeps it unmerged: $y = W_{4\text{bit}}(x) + \frac{\alpha}{r} B A x$ computed as a parallel delta, the 4-bit base bytes untouched. The alternative, merging the delta into the dequantized weight and re-quantizing to 4-bit, fails on arithmetic rather than implementation: LoRA deltas (RMS 10^{-3}) sit below the 4-bit group step, so requantization erases most of the weight-level delta (34% of the effect survives on an attention projection, 2% on a dense projection), and at the logits level 18% survives.² The unmerged path costs 42 MB of adapter weights and no measurable decode time, and it is strictly more faithful. Adapters therefore ship as files beside the checkpoint: one read-only base serves every adapter generation, and retraining a tier means swapping two files.

4 Training

The recipe has three phases, and all of them run on the dequantized bf16 reconstruction of the 4-bit deployment checkpoint (we train on what is served), with the base frozen throughout.

Phase 1: distill the heads. The prerouter heads are the only trained parameters. The loss imitates the next layer’s true router, so the heads learn to predict routing rather than to fit text.

Phase 2: SFT on the student path. The LoRA is attached to attention, linear-attention, and shared-expert projections, but not to routed experts, whose weights stream and must stay replaceable, and training runs the full forward with prerouter routing active (the “student path”) over roughly two million rows of teacher-generated text. This is the phase that makes the approximation usable: in our runs, distillation-only checkpoints with student routing produce repetitive, collapsed text, while the same heads plus SFT produce coherent output at identical speed and memory. The order was not negotiable in those runs: heads first (the SFT signal otherwise drowns the tiny head

²Retention at the logits level is measured as $1 - \|U - M\| / \|U - B\|$, which is 0.18 for the merged model, where U = unmerged, M = merged, B = base, first-token logits on a fixed prompt; weight-level retention is $\Sigma(\text{actual} \cdot \delta) / \Sigma \delta^2$ per target.

gradients), SFT second, on-policy distillation last. Chasing router agreement harder is the wrong objective: cross-token prediction from the previous token’s hidden state is information-limited, and what matters is whether the *language model* produces good text under the student routing.

Phase 3: on-policy distillation. Phase 2 trains on text the teacher wrote; Phase 3 trains on the student’s own generations. The Phase-2 checkpoint generates, the original fp16 base scores those tokens as teacher, and the gradient is taken through the served path on the same trainable surface as Phase 2. The objective is reverse KL (mode-seeking, so the student is never asked to cover the teacher’s entire support [1, 14]), applied to the teacher’s top- k tokens, with the mass outside the top- k set carried by a tail term rather than dropped [4, 16]. Phase 3 converges on one tenth of the SFT corpus (roughly 200k rows against the ≈ 2 M of Phase 2).

Iteration economics and the released width. Retraining a tier for a different routing width means swapping adapter files and nothing else. That is what let us release the 35B tier at $K=4$ rather than the base model’s $K=8$: on the 16 GB machine of §5.3, in a same-session A/B at one cache budget with the same weights at both widths, narrowing from $K=8$ to $K=4$ nearly doubles decode (3.3 to 6.4 tok/s, Table 3) and lowers peak active memory, while the retrained tier holds the quality of Table 2. Width is the one knob here that moves speed and memory together, and the frozen base is what makes turning it a file swap rather than a training campaign.

5 Evaluation

5.1 Setup

Quality runs use OpenCompass on a compute server under identical settings for EDGE0 (int4 + adapters + prerouter routing) and the original fp16 bases; they involve no timing. All throughput and memory measurements are single-device: tier-profile numbers (Table 1) on a Mac mini M4 Pro 24GB, the prerouter A/B (Fig. 3) on a MacBook M2 with 16 GB of unified memory holding the release directory (18.4 GiB: a 19.5 GB int4 base plus 0.2 GB of adapter and prerouter heads), a machine on which the weights do not fit, so every step faults experts back in from the SSD, with mlx 0.30.4–0.30.6 spanning the campaigns [2]. Speed comparisons use same-session alternating A/B with page-cache warmup and matched cache budgets on both arms: each arm runs in its own process, the arm order rotates every round, both arms replay the same sampled token sequence, and every number we report is a median over repeated runs. Single-shot benchmarks on this hardware carry $\pm 40\%$ run-to-run spread, and up to $2.3\times$ across sessions on the 16 GB machine, which is why no cross-session number is used as evidence anywhere in this section.

5.2 Quality

Table 2 and Figure 2 give the picture. The pipeline recovers most of the joint int4-plus-routing-replacement loss: mean per-benchmark gaps of 3.9 points on the 35B tier and 2.8 on the 8B, close enough that we treat the two served tiers as quality-matched to their fp16 bases.

5.3 The Advantage of the Prerouter

The experiments in this section run on a MacBook M2 with 16 GB of unified memory serving an 18.4 GiB checkpoint, on a machine where the weights do not fit. No artificial memory limit is imposed (no mlock, no wired pages, no cgroup cap, no page-cache purge), so the only constraint

Table 1: The two public release tiers, with the released checkpoint and adapters on a Mac mini M4 Pro 24 GB. Decode and memory are the latest paired measurements (each arm in its own process, warm, arm order rotated, medians over 3 rounds); memory is the MLX allocator peak at short contexts, since expert weights stream from SSD and only the decoder’s expert cache is resident. Checkpoint on disk is the int4 base; the adapter and prerouter heads add 0.2 GB. The 35B row is the $K=4$ production profile of §3.2; the 8B row is its low-memory release profile, with a shared LRU of 64 expert bundles (≈ 1.5 GiB) separate from the per-layer staged slots; enlarging that LRU to 1024 bundles costs ≈ 0.9 GiB of the budget and lifts decode to 31.8 tok/s. Prefill warm is measured within one process on a 3.1k-token prompt, cold is the first request after process start. Disabling the prerouter on the 35B tier measures 19.9 tok/s.

	edge0-35b	edge0-8b
Base model	Qwen3.6-35B-A3B [24]	Ling 3.0 tiny (MLA+MoE) [22, 29]
Layers $\times$ experts	$40 \times 256 + \text{shared}$	$24 \times 128 + \text{shared}$
Active params / token	$\approx 3\text{B}$	$\approx 1.2\text{B}$
Routing width K	4	8
Routing family	softmax-topk	sigmoid-group (8×4 , $\times 2.5$)
Quantization	int4 affine g64	int4 affine g64
Checkpoint on disk	19.5 GB	4.5 GB
Prerouter heads	33	16
LoRA	$r=16$, $\alpha=32$	$r=16$, $\alpha=32$
Decode (tok/s)	20.4	28.0
Prefill cold / warm (tok/s)	113 / 140	500 / 1102
Peak active memory	2.9 GiB	1.5 GiB

is physical memory itself and every configuration here faults experts back in from the SSD. What must fit for the engine to run *at all* is the unreclaimable MLX allocation, and the prerouter raises it: 1.72, 1.73, and 1.79 GiB for on-demand streaming against 2.33, 2.60, and 3.22 GiB ($K=2, 4, 8$). Process RSS—which counts the file-backed expert pages as well as the allocator—peaks at 5.29, 6.12, and 6.18 GiB for on-demand streaming against 4.91, 5.61, and 6.39 GiB with the prerouter. The rest of the machine is page cache, which is reclaimable and gets displaced by residency.

The advantage is the load time moved off the critical path. With every staged layer prefetching, the prerouter issues the same reads earlier—16% more bytes at $K=4$ (Table 4). On this machine the time the main thread spends blocked on expert loads falls from 154.9 to 46.5 ms per step at $K=2$, from 244.0 to 101.9 ms at $K=4$, and from 575.0 to 211.6 ms at $K=8$. The blocked time is summed over layers, and different layers’ loads overlap, so at $K=8$ the on-demand sum (575.0 ms) is slightly larger than the step it belongs to (559.1 ms). No resource is saturated while this happens: the disk read is at most 12% of the step even at its widest, the process uses about one core of eight, and the GPU runs at 35–41%. The cost being removed is serialized load latency, and removing it is what the predictor buys (Figure 3a).

Conservation, not prediction quality. Both arms read within a few percent of the same bytes per step at $K=8$ (125.1 against 126.9 MiB) and at $K=2$ (29.5 against 30.4 MiB); only at $K=4$ does the prerouter arm read more, 58.9 against 50.9 MiB, or 16% (Table 4). Moving a read off the critical path can hide it, but it cannot delete it: the advantage comes from the cold-read time that is exposed on that path today. With every staged layer prefetching, decode reaches 8.6, 6.4, and 3.3 tok/s against 4.8, 3.5, and 1.8 tok/s on demand: +80%, +82%, and +84% (Table 3). How much




Figure 2: EDGE0 (int4 + prerouter routing + recovery LoRA) vs fp16 base models, OpenCompass, identical settings. Mean per-benchmark gap 3.9 (35b) and 2.8 (8b) points.

Table 2: Quality vs fp16 base models (max 100, OpenCompass).

Benchmark	edge0-35b (int4)	Qwen3.6 (fp16)	edge0-8b (int4)	Ling 3.0 tiny (fp16)
AIME 2026	86.6	92.7	63.3	73.3
HumanEval	90.9	95.1	91.5	92.7
GPQA-Diamond	79.8	81.8	70.7	71.2
MMLU-Pro	81.0	84.6	70.1	65.8
IFBench	57.9	61.7	53.9	60.6
Average	79.2	83.2	69.9	72.7

there is to win is a property of the storage tier, not of the head: the head only has to supply the right set.

Realizing it takes reuse. Adjacent tokens agree on only about a quarter of a layer’s expert set in our traces, so most prefetched experts are never read again, and a load that is issued and not used buys nothing. The storage tier sets how much is available; reuse sets how much of it the prediction collects.

The price is page cache, and cold pages are a granularity problem. Prefetching requires the predicted experts to be resident before they are needed, and resident MLX memory is not reclaimable: at $K=8$ the prerouter arm holds 1.43 GiB more of it while the page cache loses 1.15 GiB, four fifths of what the residency takes. The same cold pages then arrive in fewer, larger loads: 1.40 MiB per load at 90% cold, against 0.32 MiB per load at 20% cold for on-demand streaming (Figure 3b, Table 4). The cost of a load is $1.17 \text{ ms} + 1.33 \text{ ms} \times (\text{cold fraction})$, fitted on the $K=8$ pair and reproducing every measured per-load cost within 0.13 ms (the prerouter points within 0.07 ms). Prefetching is cheap per load and expensive per byte moved early.




Figure 3: The advantage of the prerouter, on a 16 GB MacBook M2 with an 18.4 GiB checkpoint that does not fit. (a) Decode throughput: pure on-demand streaming against every staged layer prefetching its next token’s experts. (b) The same cold pages read in fewer, larger loads, with the cold fraction of each load printed above its bar; per-load cost follows $1.17 \text{ ms} + 1.33 \text{ ms} \times \text{cold}$. Same-session rotated A/B, both arms replaying the same sampled token sequence, 3-round medians, 3/3 rounds agreeing.

Table 3: The advantage, on the same machine and in the same session. Decode throughput with on-demand streaming against every staged layer prefetching one token ahead.

K	on-demand	prerouter	gain
2	4.8	8.6	+80%
4	3.5	6.4	+82%
8	1.8	3.3	+84%

Two levers not yet pulled. Since the price is unreclaimable residency, the cheapest remaining gain would be to stop paying for memory the loader does not need: the stager in these runs keeps both a bundle and a stacked tensor view of the same weights, and dropping one copy would return about 0.45 GiB at $K=8$. The second is the fill itself: with incremental stacking it runs synchronously on the main thread here, which at $K=8$ costs 62.5 ms per step, so moving it off that thread would remove a cost that never needed to be waited for. Neither change is in the configuration measured above.

One mechanism, not two. The prerouter supplies staged decode with a set that is correct *by definition* (it is the routing), so the staged slots map exactly and nothing is dropped. Staging alone is not a deployable configuration: without a correct set source it either drops experts and the output degrades, or pins a hot set that churns. Prediction chooses, staging loads.

5.4 Memory: What the Machine Actually Pays

Peak active memory is 2.9 GiB for the 35B tier against a 19.5 GB checkpoint (Table 1); the remainder is disk. The same weights fully resident need 18.2 GiB and decode at 3.9 tok/s: 18.2 GiB of unreclaimable weights leave too little of the 24 GB for the KV cache and the operating system, so

Table 4: What the price is made of (same session, same arms). Both arms move within a few percent of the same bytes at $K=2$ and $K=8$ (29.5 against 30.4 and 125.1 against 126.9 MiB, prerouter against on-demand); at $K=4$ the prerouter arm reads 16% more (58.9 against 50.9 MiB). The prerouter reads them in fewer, larger, colder loads.

K	arm	loads/step	MiB/step	MiB/load	cold	ms/load
2	on-demand	99.3	30.4	0.31	20%	1.56
2	prerouter	19.3	29.5	1.53	98%	2.41
4	on-demand	189.4	50.9	0.27	17%	1.29
4	prerouter	44.7	58.9	1.32	84%	2.28
8	on-demand	398.6	126.9	0.32	20%	1.44
8	prerouter	89.5	125.1	1.40	90%	2.36

the system pages. EDGE0 occupies one-sixth to one-seventh of what the resident server holds and decodes at 20.4 tok/s (Table 1), five times faster.

Prefill behaves as the mirror image: cold first-request prefill pays the SSD fault-ins (35B tier: 113/140 tok/s cold/warm on a 3.1k-token prompt; 8B tier: 500/1102), and whole-layer loading plus the page cache makes warm prefill compute-bound. Decode, not prefill, is the constrained phase, and every mechanism here targets it.

6 Limitations

EDGE0 serves one request at a time, FIFO-serialized. Concurrency belongs to a serving layer, not the engine, and batching changes the expert working set in ways our per-request profiles do not model.

The streaming engine’s decode is CPU-side, not storage-side: 44 ms per step of graph building in the 40-layer forward is the floor, and no storage-side optimization moves it. Closing that gap needs kernel-level graph amortization or a smaller model (Appendix B).

Prerouter gains scale with the waiting they remove: they shrink on hot caches and fast storage, and they are bounded by reuse: adjacent tokens agree on only about a quarter of a layer’s experts, so a prefetched expert is often read once and paid for twice. Its price is unreclaimable residency taken from the page cache, which is exactly what makes the mechanism a win where memory is tight and the storage tier is slow, the regime EDGE0 was built for.

Quality loss concentrates in long-chain reasoning: 6.1 points on AIME for the 35B tier and 10.0 for the 8B tier; every other 8B benchmark is within 6.7 points, and MMLU-Pro is 4.3 points in the student’s favour. The recovery LoRA recovers most of the pipeline’s loss everywhere else, and reasoning is where int4 plus routing replacement remains visible.

The MLX backend is the one implementation; the backend facade is the abstraction, and the CUDA slot is architecture, not code.

7 Conclusion

The weights half of the memory wall is a storage placement decision, and everyone made it the same way. EDGE0 makes the other decision: experts live on SSD, a trained prerouter predicts routing one token ahead so loads hide under compute, prediction is the routing so nothing is dropped, and a distilled, unmerged LoRA pays the quality bill at 4-bit. The result is a 35B-class MoE served

from a 24 GB consumer desktop at 20 tok/s inside 3 GiB of memory, an 8B hybrid at 28 tok/s inside 1.5 GiB, both within a few points of their fp16 teachers, with every component (framework, checkpoints, adapters) open. The machine on your desk is already big enough; the weights just needed somewhere to live.

Artifacts. Framework: <https://github.com/Edge0-AI/Edge0>. Models: <https://huggingface.co/Edge0/Edge0-35B-A3B-preview> and <https://huggingface.co/Edge0/Edge0-8B-A1B-preview>. Apache-2.0.

References

- [1] Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu Geist, et al. On-policy distillation of language models: Learning from self-generated mistakes. *arXiv preprint arXiv:2306.13649*, 2023.
- [2] Apple. MLX: An array framework for apple silicon. <https://github.com/ml-explore/mlx>, 2023.
- [3] Runyuan Cai, Yiming Wang, Yu Lin, and Xiaodong Zeng. Tiny-engram: Trigger-indexed concept tables for generative vision. *arXiv preprint arXiv:2605.20309*, 2026.
- [4] Sayantan Dasgupta, Trevor Cohn, and Timothy Baldwin. Don’t ignore the tail: Decoupling top- k probabilities for efficient language model distillation. *arXiv preprint arXiv:2602.20816*, 2026.
- [5] DeepSeek-AI. DeepSeek-V2: A strong, economical, and efficient mixture-of-experts language model. *arXiv preprint arXiv:2405.04434*, 2024.
- [6] DeepSeek-AI. DeepSeek-V3 technical report. *arXiv preprint arXiv:2412.19437*, 2024.
- [7] DeepSeek-AI. DeepSeek-V3.2-Exp: Boosting long-context efficiency with DeepSeek sparse attention. <https://github.com/deepseek-ai/DeepSeek-V3.2-Exp>, 2025. Introduces DeepSeek Sparse Attention (DSA).
- [8] Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. QLoRA: Efficient finetuning of quantized LLMs. *arXiv preprint arXiv:2305.14314*, 2023.
- [9] Artyom Eliseev and Denis Mazur. Fast inference of mixture-of-experts language models with offloading. *arXiv preprint arXiv:2312.17238*, 2023.
- [10] William Fedus, Barret Zoph, and Noam Shazeer. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *arXiv preprint arXiv:2101.03961*, 2021.
- [11] Elias Frantar, Saleh Ashkboos, Torsten Hofer, and Dan Alistarh. GPTQ: Accurate post-training quantization for generative pre-trained transformers. In *International Conference on Learning Representations (ICLR)*, 2023.
- [12] Georgi Gerganov and contributors. llama.cpp. <https://github.com/ggml-org/llama.cpp>, 2023.
- [13] Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. In *Conference on Language Modeling (COLM)*, 2024.

- [14] Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. MiniLLM: On-policy distillation of large language models. *arXiv preprint arXiv:2306.08543*, 2024.
- [15] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations (ICLR)*, 2022.
- [16] Huipeng Huang and Hongxin Wei. Tail-aware top- k on-policy distillation. *arXiv preprint arXiv:2608.14728*, 2026.
- [17] Ranggi Hwang, Jianyu Wei, Shijie Cao, Changho Hwang, Xiaohu Tang, Ting Cao, and Mao Yang. Pre-gated MoE: An algorithm-system co-design for fast and scalable mixture-of-expert inference. In *Proceedings of the 51st IEEE/ACM International Symposium on Computer Architecture (ISCA)*, 2024. URL <https://arxiv.org/abs/2308.12066>. arXiv:2308.12066.
- [18] Albert Q. Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, et al. Mixtral of experts. *arXiv preprint arXiv:2401.04088*, 2024.
- [19] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. *arXiv preprint arXiv:2309.06180*, 2023.
- [20] Dmitry Lepikhin, Hyoungho Lee, Yuanzhong Xu, Dehao Chen, Orhan Firat, Yanping Huang, Maxim Krikun, Noam Shazeer, and Zhifeng Chen. GShard: Scaling giant models with conditional computation and automatic sharding. *arXiv preprint arXiv:2006.16668*, 2020.
- [21] Ji Lin, Jiaming Tang, Haotian Tang, Shang Yang, Wei-Ming Chen, Wei-Chen Wang, Guangxuan Xiao, Xingyu Dang, Chuang Gan, and Song Han. AWQ: Activation-aware weight quantization for LLM compression and acceleration. In *Proceedings of the MLSys Conference*, 2024.
- [22] Ling Team, Ant Group. Ling-3.0-tiny model release. <https://huggingface.co/inclusionAI/Ling-3.0-tiny-base>, 2026.
- [23] Xingtai Lv, Li Sheng, Kaiyan Zhang, et al. Post-trained MoE can skip half experts via self-distillation. *arXiv preprint arXiv:2605.18643*, 2026.
- [24] Qwen Team. Qwen3.6-35B-A3B model release. <https://huggingface.co/Qwen/Qwen3.6-35B-A3B>, 2026.
- [25] Samyam Rajbhandari, Conglong Li, Zhewei Yao, Minjia Zhang, Reza Yazdani Aminabadi, Ammar Ahmad Awan, Jeff Rasley, and Yuxiong He. DeepSpeed-MoE: Advancing mixture-of-experts inference and training to power next-generation AI scale. *arXiv preprint arXiv:2201.05596*, 2022.
- [26] Noam Shazeer, Azalia Mirhoseini, Krzysztof Maziarczyk, Andy Davis, Quoc Le, Geoffrey Hinton, and Jeff Dean. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. In *International Conference on Learning Representations (ICLR)*, 2017.
- [27] Ying Sheng, Lianmin Zheng, Binhang Yuan, Zhuohan Li, Max Ryabinin, Daniel Y. Fu, Zhiqiang Xie, Beidi Chen, Clark Barrett, Joseph E. Gonzalez, Percy Liang, Christopher Ré, Ion Stoica, and Ce Zhang. FlexGen: High-throughput generative inference of large language models with a single GPU. *arXiv preprint arXiv:2303.06865*, 2023.

- [28] Yixin Song, Zeyu Mi, Haotong Xie, and Haibo Chen. PowerInfer: Fast large language model serving with a consumer-grade GPU. *arXiv preprint arXiv:2312.12456*, 2023.
- [29] Ling Team et al. Every FLOP counts: Scaling a 300b mixture-of-expert LING LLM without premium GPUs. *arXiv preprint arXiv:2503.05139*, 2025.
- [30] Yiming Wang, Yu Lin, Xiaodong Zeng, and Guannan Zhang. MultiLoRA: Democratizing LoRA for better multi-task learning. *arXiv preprint arXiv:2311.11501*, 2023.
- [31] Yiming Wang, Yu Lin, Xiaodong Zeng, and Guannan Zhang. PrivateLoRA for efficient privacy preserving LLM. *arXiv preprint arXiv:2311.14030*, 2023.
- [32] W. A. Wulf and S. A. McKee. Hitting the memory wall: Implications of the obvious. *ACM SIGARCH Computer Architecture News*, 23(1):20–24, 1995.
- [33] Leyang Xue, Yao Fu, Zhan Lu, Luo Mai, and Mahesh Marina. MoE-Infinity: Efficient MoE inference on personal machines with sparsity-aware expert cache. *arXiv preprint arXiv:2401.14361*, 2024.

A Routing Math

Both functions are extracted verbatim from the vendored model implementations and shared by resident, streaming, and prerouter paths; parity tests pin them bit-identical.

Softmax-topk (K , renormalize).

$$\begin{aligned} g &= \text{softmax}(\ell) & \text{inds} &= \text{top-k}(g, K) \\ w &= g[\text{inds}] / \sum_{\text{inds}} g \end{aligned}$$

Sigmoid-group (K , G of n groups, scale s).

$$\begin{aligned} \sigma &= \text{sigmoid}(\ell) \\ \text{group score}_i &= \text{top-2 sum of } \sigma \text{ in group } i, \quad \text{keep top } G \text{ groups, mask others to } -\infty \\ \text{inds} &= \text{top-k}(\sigma \text{ within survivors}, K) \\ w &= s \cdot \sigma[\text{inds}] / (\sum \sigma[\text{inds}] + 10^{-20}) \end{aligned}$$

Selection uses the biased score; weights use the raw sigmoid, the DeepSeek-style distinction [6]. The 8B tier runs $n=8$, $G=4$, $s=2.5$, $K=8$.

B The Assembly Tax and Incremental Stacks

The oracle bound (§3.1) locates the residual gap in per-step tensor assembly. In a separate A/B against the same pipeline without `incr_stack`, replacing per-step stacks with persistent sticky-slot tensors updated in place contributes a further +34% decode. In a second same-session A/B on the $K=8$ configuration—native routing against the complete shipped pipeline—decode rises from 6.8 to 12.5 tok/s. Both pairs come from the earlier campaign rather than the same-session sweep of §5.3, and from a different configuration, so their rates are not comparable with Table 3; they price assembly, not prefetching: with `incr_stack` the pipeline rewrites only the slots that changed, and what the prerouter does *not* remove is the per-step assembly that remains.